

A Geometric Multigrid Preconditioner for Shifted Boundary Method

Michał Wichrowski*, Ajay Ajith †

November 13, 2025

1 Introduction

2 Method formulation

MW: Introduction needs to be rewritten.

We consider the Poisson problem as a model problem:

$$-\Delta u = f \quad \text{in } \Omega, \quad (1)$$

$$u = g \quad \text{on } \Gamma, \quad (2)$$

where $\Omega \subset \mathbb{R}^d$ ($d = 2, 3$) is a domain with boundary $\Gamma = \partial\Omega$ as depicted in Figure 1, f is a given source term, and g is the prescribed Dirichlet boundary condition.

To solve this problem numerically, we first formulate it in a weak sense. We seek a solution u in an appropriate function space, $H^1(\Omega)$, which consists of functions that are square-integrable and whose first derivatives are also square-integrable. Multiplying the equation by a test function $v \in H^1(\Omega)$ and integrating over Ω , we obtain:

$$\int_{\Omega} \nabla u \cdot \nabla v \, dx - \int_{\Gamma} (\nabla u \cdot \mathbf{n}) v \, ds = \int_{\Omega} f v \, dx.$$

We next introduce a triangulation $\mathcal{T}_h$ consisting of quadrilateral (2D) or hexahedral (3D) elements of size h , and define a finite element space $\mathbb{V}_h \subset H^1(\Omega)$ using Lagrange polynomial elements of degree p . In classical finite element methods, the mesh conforms to the boundary (unlike the background mesh approach illustrated in Figure 1), and test functions v typically vanish on Γ to strongly enforce the Dirichlet condition $u = g$. When function spaces do not necessarily satisfy essential boundary conditions strongly, boundary conditions must be enforced weakly through additional integral terms on Γ .

Nitsche's method provides a way to weakly impose Dirichlet boundary conditions within a variational formulation without requiring the function space to satisfy the boundary conditions. It modifies the bilinear form by adding terms on the boundary Γ . The standard Nitsche formulation for the Poisson problem with Dirichlet boundary conditions $u = g$ on Γ seeks $u_h \in \mathbb{V}_h$ such that for all $v_h \in \mathbb{V}_h$:

$$\begin{aligned} \int_{\Omega} \nabla u_h \cdot \nabla v_h \, dx - \int_{\Gamma} (\nabla u_h \cdot \mathbf{n}) v_h \, ds - \alpha \int_{\Gamma} (\nabla v_h \cdot \mathbf{n}) u_h \, ds + \int_{\Gamma} \sigma u_h v_h \, ds = \\ = \int_{\Omega} f v_h \, dx - \int_{\Gamma} (\nabla v_h \cdot \mathbf{n}) g \, ds + \int_{\Gamma} \sigma_{\Gamma} g v_h \, ds. \end{aligned}$$

*Heidelberg University, Germany, mt.wichrowsk@uw.edu.pl

†Heidelberg University, Germany

Here, σ_Γ is a penalty parameter, typically chosen as $\sigma_\Gamma = \mathcal{O}(p^2 h^{-1})$ for mesh size h , and $\mathbf{n}$ is the outward unit normal vector to Γ . The choice of parameter $\alpha = 1$ leads to a symmetric formulation, while $\alpha = -1$ results in a non-symmetric formulation in which the penalty term can be skipped [2]




Figure 1: Schematic illustrating the background mesh, interior cells (green), the surrogate boundary $\tilde{\Gamma}$ (thick blue line) along the upper boundary of the interior cells, and the true boundary $\partial\Omega$ (red).

2.1 Shifted Boundary Method

In many applications, the domain Ω may have a complex geometry, making the generation of body-fitted meshes challenging. Non-body-fitted (unfitted) methods address this challenge by employing a background mesh $\mathcal{T}_h$ that does not conform to the boundary of Ω . The domain Ω is embedded within this background mesh, and cells of $\mathcal{T}_h$ are classified as active based on their intersection with Ω . The computational domain $\tilde{\Omega}$ is defined as the union of these active cells. In the original SBM formulation [5], only cells strictly contained in Ω were included, leading to the surrogate domain boundary depicted by the solid blue line in Figure 1. Consequently, the surrogate domain $\tilde{\Omega}$ is generally a subset of Ω , and its boundary $\tilde{\Gamma}$ does not coincide with the true boundary Γ .

Later extensions include intersected cells [8] with a volume fraction inside Ω greater than a threshold $\lambda \in [0, 1]$. In Figure 1, this corresponds to including additional cells as indicated by the dashed blue line. This approach reduces the distance between true and surrogate boundaries, but can lead to ill-conditioning.

To impose boundary conditions on $\tilde{\Gamma}$, the Dirichlet condition prescribed on the true boundary Γ is extrapolated to the surrogate boundary. This is typically accomplished by projecting points from $\tilde{\Gamma}$ onto the true boundary Γ along a suitable direction (e.g., the outward normal to $\tilde{\Gamma}$) and using a Taylor expansion to approximate the boundary values. This is accomplished via an *extension operator* $\mathcal{E}$, which maps functions defined on $\tilde{\Omega}$ to the surrogate boundary $\tilde{\Gamma}$.

We assume that the Dirichlet boundary condition g is given as a restriction of a function u^* defined on the entire domain Ω to the boundary Γ . While the choice of this function u^* (as an extension of g from Γ into Ω) is not unique, we take u^* to be equal to the solution u in the surrogate domain $\tilde{\Omega}$. For each point $\tilde{\mathbf{x}} \in \tilde{\Gamma}$, let $\mathbf{x} \in \Gamma$ be its closest point projection onto the true boundary, and let $\mathbf{d} = \mathbf{x} - \tilde{\mathbf{x}}$ be the shift vector. The function u is extended from Γ to $\tilde{\Gamma}$ using a Taylor expansion:

$$\mathcal{E}u^*(\tilde{\mathbf{x}}) = u^*(\mathbf{x}) - \mathbf{d} \cdot \nabla u(\tilde{\mathbf{x}}) + \dots$$

where $\mathcal{E}u^*$ denotes the extrapolated boundary condition. The function u^* is assumed to be smooth in a neighborhood of $\tilde{\Gamma}$, which allows for the Taylor expansion to be valid. By substituting it into the weak formulation, we obtain a variational formulation for the shifted boundary problem with the extrapolated boundary condition on $\tilde{\Gamma}$ enforced in a Nitsche-like manner; the

MW: This section needs to be altered, for not it is a copy-paste from the previous paper.

weak formulation seeks $u_h \in V_h$ such that for all $v_h \in V_h$,

$$\begin{aligned} \int_{\tilde{\Omega}} \nabla u_h \cdot \nabla v_h \, dx - \int_{\tilde{\Gamma}} (\nabla u_h \cdot \tilde{\mathbf{n}}) v_h \, ds - \alpha \int_{\tilde{\Gamma}} (\nabla v_h \cdot \tilde{\mathbf{n}}) \mathcal{E} u_h \, ds + \int_{\tilde{\Gamma}} \sigma \mathcal{E} u_h v_h \, ds = \\ = \int_{\tilde{\Omega}} f v_h \, dx - \int_{\tilde{\Gamma}} (\nabla v_h \cdot \tilde{\mathbf{n}}) g^E \, ds + \int_{\tilde{\Gamma}} \sigma g v_h \, ds. \end{aligned} \quad (3)$$

where $\tilde{\mathbf{n}}$ is the outward normal to $\tilde{\Gamma}$ and σ is a penalty parameter.

The choice of the stabilization term, particularly the parameter α , significantly influences the spectral properties of the resulting system matrix. As detailed in our previous work on a Discontinuous Galerkin (DG) based SBM multigrid preconditioner [7], this can be observed even in a simple 1D single-cell problem. For instance, a penalty-free formulation ($\sigma = 0$, $\alpha = -1$) can lead to complex eigenvalues, especially when the true boundary lies inside the surrogate domain.

In the context of DG methods, it was found that the choice of stabilization did not have a dramatic impact on the overall multigrid performance, as the cell-based nature of the DG smoother effectively handled local issues [7]. However, in the DG formulation the cellwise smoother experienced significant challenges for $p = 3$. For continuous finite elements, where our smoothers operate on larger, overlapping patches of elements (i.e. vertex or edge patches), the choice of stabilization becomes far more critical. The non-standard terms introduced by SBM are not well-contained within these patches, and an inappropriate stabilization can introduce spectral properties that standard smoothers cannot handle effectively, a problem that is exacerbated for higher-order elements.

The SBM weak formulation with symmetrized stabilization seeks $u_h \in V_h$ such that for all $v_h \in V_h$,

$$\begin{aligned} \int_{\tilde{\Omega}} \nabla u_h \cdot \nabla v_h \, dx - \int_{\tilde{\Gamma}} (\nabla u_h \cdot \tilde{\mathbf{n}}) v_h \, ds - \int_{\tilde{\Gamma}} (\nabla v_h \cdot \tilde{\mathbf{n}}) \mathcal{E} u_h \, ds + \int_{\tilde{\Gamma}} \sigma_{\Gamma} \mathcal{E} u_h \mathcal{E} v_h \, ds = \\ = \int_{\tilde{\Omega}} f v_h \, dx - \int_{\tilde{\Gamma}} (\nabla v_h \cdot \tilde{\mathbf{n}}) g^E \, ds + \int_{\tilde{\Gamma}} \sigma_{\Gamma} g v_h \, ds. \end{aligned} \quad (4)$$

where $\tilde{\mathbf{n}}$ is the outward normal to $\tilde{\Gamma}$ and σ_{Γ} is a penalty parameter. We notice that our discrete solution u_h is a piecewise polynomial function defined on the background mesh $\mathcal{T}_h$, hence its Taylor expansion can be computed directly by evaluating the function values at the points on the true boundary Γ . This allows us to avoid computation of higher-order derivatives of u_h . Note that the resulting form is not symmetric due to the presence of the extrapolated boundary condition.

3 Multigrid Preconditioner

The SBM formulation described above leads to a large, sparse linear system. Due to the lack of symmetry and possible indefiniteness of the resulting matrix, we employ a Krylov subspace method, specifically GMRES, for its solution. The convergence of GMRES is sensitive to the condition number of the system matrix, which grows with both the number of elements in the mesh and the polynomial degree. To accelerate the solution, we employ a multigrid preconditioner [4]. Our Cartesian background mesh naturally facilitates the construction of a nested hierarchy of meshes $\{\mathcal{T}_{\ell}\}_{\ell=0}^L$, where ℓ denotes the level and L is the finest level. On each mesh $\mathcal{T}_{\ell}$, we define a finite element space $\mathbb{V}_{\ell}$ consisting of continuous piecewise polynomials of a fixed degree p .

This represents an important difference from the work on a Discontinuous Galerkin SBM preconditioner [7], where an hp-multigrid strategy with lower polynomial degrees on coarser levels was necessary to achieve good performance. In the present continuous Galerkin context,

we find that a simpler h-multigrid approach is enough. We note that while the background meshes are nested, the sets of active cells on different levels (determined by the fraction λ of their volume inside the domain) are not necessarily nested. This lack of geometric nestedness can affect the efficiency of standard inter-grid transfer operators.

With this hierarchy of meshes and spaces established, the multigrid method is defined through three crucial components: a smoother, which reduces high-frequency error components on each level; transfer operators, which move information between different resolution levels; and a coarse-grid solver. For preconditioning, we use a single multigrid V-cycle.

3.1 Smoother

We first consider Richardson iterations with preconditioners P . Given the current approximation u_h and the right-hand side f_h , a smoothing step updates the approximation:

$$u_h^{k+1} = u_h^k + P(f_h - A_h u_h^k).$$

We decompose the space $\mathbb{V}_\ell$ into a sum of subspaces $\mathbb{V}_i$, i.e., $\mathbb{V}_\ell = \sum_{i=1}^N \mathbb{V}_i$. Let A_i be the restriction of A to the subspace $\mathbb{V}_i$. Then, the additive subspace correction preconditioner is defined as:

$$P = \omega \sum_{i=1}^N A_i^{-1}$$

where ω is a relaxation parameter. Alternatively, the successive subspace correction method is a subspace correction method where the subspaces $\mathbb{V}_i$ are visited in a sequential manner. The procedure can be performed by applying one Richardson iteration with the preconditioner P_i for each subspace $\mathbb{V}_i$. Since in each step only one subspace is corrected, the residual only changes locally and an efficient implementation is possible. The preconditioner updates the solution by sequentially applying corrections for each subspace. For each subspace $\mathbb{V}_i$, a correction is computed and applied:

$$u_h \leftarrow u_h + R_i^T A_i^{-1} R_i (f_h - A_h u_h).$$

This process is repeated for all subspaces $i = 1, \dots, N$. If the subspaces are chosen as one-dimensional spaces spanned by the basis functions, then the preconditioner is equivalent to either Jacobi (additive) or Gauss-Seidel (successive).

Viewing the smoother as a subspace correction method provides a useful framework. The effectiveness of such a smoother depends critically on the choice of subspaces and the local problems solved on them. For continuous finite element methods, simple choices like one-dimensional subspaces corresponding to individual degrees of freedom (leading to Jacobi or Gauss-Seidel smoothers) are ineffective for SBM. These methods implicitly enforce a zero correction on the boundary of the single-vertex *patch*, which conflicts with the non-standard boundary terms introduced by SBM, a fact confirmed by our preliminary experiments.

This suggests that the subspaces must be large enough to capture the local behavior of the SBM formulation. We therefore construct local problems on larger subspaces, defined by patches of elements centered around a vertex, as illustrated in Figure 2. By solving a local problem on this patch that includes the relevant SBM boundary terms where applicable, we can compute a more effective correction. This approach aims to make the local problem on the patch a more consistent approximation of the global problem, leading to a more robust smoother.

3.2 Full-Residual Shy Patches

The idea behind our proposed smoother is to construct local problems on patches of cells that are consistent with the global problem, particularly in handling the non-standard terms introduced




Figure 2: Left: Boundaries of a vertex patch. Interior boundaries shown in light gray, and the thick blue segments indicate the exterior faces where SBM boundary conditions are applied. Right: the same vertex patch with a gray frame highlighting which degrees of freedom form the patch subspace; three other patches are indicated with gray dashed outlines, and the central vertices are marked by gray squares.

by the SBM. We start by constructing vertex patches, where a patch consists of all active cells connected to a given vertex. The subspace for the smoother correction is then defined by the degrees of freedom (DoFs) supported by the cells within this patch.

Next we ensure that the local problem solved on the patch is a faithful representation of the global problem. To achieve this, we only include DoFs for which the complete residual can be evaluated using information solely from the cells within the patch. This means we exclude any DoFs whose support extends to cells outside the current patch. By doing so, we guarantee that the local correction is computed using the true global residual, avoiding inconsistencies that would arise from incomplete information, especially near the surrogate boundary. Figure 2 illustrates this concept for a few patches.

We further refine this patch selection with a concept we call *shyness*. Much like a shy person at a party needs a circle of friends to feel confident enough to join the dance, our vertices are considered shy. A vertex will only form the center of a patch if it is surrounded by a sufficient number of *friends* — in this case, active cells. We define a shyness threshold as the minimum number of active cells required around a vertex to consider it for patch construction. If a vertex does not meet this threshold, it is deemed too isolated, and no patch is formed around it. This strategy prevents the creation of small local problems, which reduces the computational cost of the smoother and avoids potential issues with ill-conditioning that can arise from patches with too few active cells.

In [7], satisfactory results were obtained with 3 smoothing steps. A multi-stage smoothing strategy was proposed in [3] for CutFEM, where an initial global smoothing pass was followed by additional passes restricted to patches containing cut cells. We adopt a similar strategy here. The first smoothing step is applied to all patches in the domain. Subsequent steps are then selectively applied only to those patches that contain cells adjacent to the surrogate boundary. This focuses the computational effort of the smoother on the region where the SBM introduces non-standard terms and where the error is often most difficult to resolve. These boundary-adjacent patches are identified once during the construction of the smoother.

3.3 Transfer Operators

The transfer of information between different levels of the multigrid hierarchy is handled by prolongation and restriction operators. The prolongation operator, $I_{\ell-1}^{\ell}$, maps a function from the coarse space $\mathbb{V}_{\ell-1}$ to the fine space $\mathbb{V}_{\ell}$, while the restriction operator, $I_{\ell}^{\ell-1}$, transfers a function from the fine space to the coarse space. For standard multigrid methods on nested meshes with continuous finite elements, prolongation is typically the natural embedding of the coarse function space into the fine one, and restriction is its transpose.

A significant challenge in applying multigrid to unfitted methods like SBM is the lack of geometric nestedness of the computational domains. The set of active cells $\tilde{\Omega}_\ell$ on a given level ℓ is determined independently based on the intersection with the true domain Ω . Consequently, the active domain on the fine level, $\tilde{\Omega}_L$, is not necessarily a subset of the active domain on a coarser level, $\tilde{\Omega}_\ell$ (when viewed on the fine grid). This can lead to inconsistencies where, for example, an active fine-grid cell corresponds to a non-active coarse-grid cell.

This lack of nestedness can impair the effectiveness of standard transfer operators [7]. Information from an active region on the fine grid might be restricted to a non-active region on the coarse grid, where it is essentially discarded, breaking the flow of information required for an efficient multigrid cycle.

4 Implementation details

The numerical implementation of our multigrid solver is based on the open-source finite element library `deal.II` [1]. It provides a comprehensive framework for the implementation of finite element methods, including mesh handling, finite element spaces, assembly of linear systems, and interfaces to various linear algebra solvers and preconditioners. In this paper we build upon the implementation of the multigrid solver for DG-SBM presented in [7].

4.1 Background mesh and geometry handling

When implementing SBM on a non-body-fitted mesh, we need to handle cells intersected by the true boundary Γ . We use a level set function $\phi(\mathbf{x})$ to implicitly define the domain $\Omega = \{\mathbf{x} \mid \phi(\mathbf{x}) < 0\}$, with $\Gamma = \{\mathbf{x} \mid \phi(\mathbf{x}) = 0\}$. Cells of the background mesh are classified based on their intersection with the zero level set: cells entirely inside Ω (interior), cells entirely outside Ω (exterior), and cells intersected by Γ . Then, for the intersected cells the fraction of the cell volume inside Ω is computed, and the cell is classified as active if this fraction is greater than a threshold λ . This is handled using non-matching quadrature rules implemented in `deal.II`, which are based on the techniques described in [6].

In our approach, degrees of freedom are formally assigned to all cells of the background mesh, including those that are classified as non-active. However, for non-active cells, the local contribution to the global system matrix is set to the identity matrix, and the corresponding entries in the right-hand side are set to zero. This effectively decouples the degrees of freedom in non-active cells from the system, ensuring that the solution is only computed within the computational domain $\tilde{\Omega}$ while simplifying the global matrix assembly process by maintaining a uniform structure of degrees of freedom (DoF) across all cells.

4.2 Processing surrogate boundary and matrix assembly

The SBM requires computing the closest point projection from points on the surrogate boundary $\tilde{\Gamma}$ to the true boundary Γ . This projection is found by solving a local nonlinear optimization problem for each quadrature point on $\tilde{\Gamma}$, which minimizes the distance to the true boundary defined by the zero level set of a function $\phi(\mathbf{x})$. The full details of the formulation, which involves a Lagrange multiplier approach solved with a Newton-Raphson method, are described in [7]. To ensure sufficient smoothness for the derivative calculations required by the solver, the level set function $\phi(\mathbf{x})$ is represented using finite elements of degree 2 for $p = 1$ and degree p for $p > 1$. The search for the closest point is performed only in the interior of the cells adjacent to the surrogate boundary. While this does not guarantee that the closest point is found inside the cell, it is expected that even if the closest point is outside the cell, the extrapolation will still yield a good approximation of the boundary condition.

The extrapolation of the function values from the true boundary to the surrogate boundary required for the matrix assembly process is accomplished by evaluating the values of the basis

functions at the points on the surrogate boundary. The assembly of the cell contributions and interior faces to the system matrix and right-hand side vector is performed using the standard finite element assembly process provided by `deal.II`.

4.3 Multigrid structures

For the multigrid hierarchy, we leverage `deal.II`'s built-in capabilities for handling nested meshes and defining finite element spaces on each level. The standard L^2 projection operators provided by `deal.II` are used for the prolongation ($I_{\ell-1}^\ell$) and restriction ($I_\ell^{\ell-1}$) operators, transferring data between coarser and finer grid levels.

The assembly of the system matrix A_ℓ on each level ℓ of the multigrid hierarchy follows a procedure analogous to that on the finest level. The level set function defining the domain geometry is interpolated onto the mesh $\mathcal{T}_\ell$. Based on this interpolated level set and the chosen threshold λ , active cells for level ℓ are identified. The SBM bilinear form is then used to assemble the local contributions to A_ℓ only for these active cells. The cells deemed non-active on level ℓ are ignored during the assembly, effectively decoupling their degrees of freedom from the system on that level.

5 Numerical results

In this section, we present numerical results to evaluate the performance of the proposed SBM multigrid preconditioner for the Poisson equation. We investigate its effectiveness in terms of convergence rates, iteration counts under mesh refinement (h -refinement) and polynomial degree increase (p -refinement). To better illustrate the multigrid preconditioner's performance characteristics and enable meaningful comparison of iteration counts, we solve the system to high accuracy with a tolerance of 10^{-12} for the relative residual reduction. This stringent tolerance results in higher iteration counts, making the differences in solver performance more apparent and facilitating clearer assessment of the multigrid effectiveness. Solver failure is defined as exceeding 100 GMRES iterations without reaching this tolerance. The initial background mesh is a square (or cube in 3D) domain covering the range $[-1.01, 1.01]$ in each dimension, subdivided into 4 cells in each coordinate direction. We use a penalty parameter $\sigma = 5$ in the SBM formulation. This parameter is selected to be sufficiently large to ensure stability of the Nitsche coupling on the surrogate boundary.

The majority of our tests are conducted on a unit sphere (a unit disk in 2D, depicted on the left panel of Fig. ??). While geometrically simple, the unit sphere serves as an insightful benchmark. Any sufficiently smooth (C^1) complex boundary, when viewed at a fine enough mesh resolution, locally resembles a flat plane. The unit sphere, due to its uniform curvature, presents a comprehensive range of intersection angles between the true boundary and the background mesh cells. Furthermore, the boundary intersects cells at various locations relative to cell centers and faces, leading to a diverse distribution of shift vector magnitudes and directions. This includes scenarios where the shift vector points from the surrogate boundary towards the interior of the true domain (which we denote as negative shifts if they oppose the outward normal of the surrogate boundary). The right panel of Figure ?? depicts the minimum and maximum shift magnitudes observed across the surrogate boundary for a typical discretization. The shift magnitudes are normalized by the cell size h ; for a unit cell, the theoretically largest possible shift magnitude in 2D is $\sqrt{2}h$, occurring if the surrogate boundary point is at a cell corner and the true boundary passes through the diagonally opposite corner.

The primary metric for evaluating the multigrid preconditioner is the number of GMRES iterations required to reduce the initial residual by a factor of 10^{-12} . If the threshold is not met within 100 iterations, the solver is considered to have failed. All experiments are conducted using the implementation described in the previous section. As our solver depends on the value

MW: don't
touch it for
now. We
might want
lambda here

of the threshold parameter λ , we will first explore the solver performance on 2D problems as they are less computationally intensive and then show some results for 3D problems with tuned λ .

6 Conclusion

Acknowledgments

The author declares the use of language models (ChatGPT, Gemini, and Claude) to improve the clarity and readability of the manuscript. All scientific content and technical claims are solely the responsibility of the author.

References

- [1] D. ARNDT, W. BANGERTH, M. BERGBAUER, B. BLAIS, M. FEHLING, R. GASSMÖLLER, T. HEISTER, L. HELTAI, M. KRONBICHLER, M. MAIER, ET AL., *The deal.ii library, version 9.7*, Journal of Numerical Mathematics, (2025).
- [2] C. E. BAUMANN AND J. T. ODEN, *A discontinuous hp finite element method for convection—diffusion problems*, Computer Methods in Applied Mechanics and Engineering, 175 (1999), pp. 311–341.
- [3] C. CUI AND G. KANSCHAT, *A multigrid method for cutfem and its implementation on gpu*, arXiv preprint arXiv:2508.11608, (2025).
- [4] W. HACKBUSCH AND W. HACKBUSCH, *The Multi-Grid Method of the Second Kind*, Multi-Grid Methods and Applications, (1985), pp. 305–353.
- [5] A. MAIN AND G. SCOVAZZI, *The shifted boundary method for embedded domain computations. Part I: Poisson and Stokes problems*, Journal of Computational Physics, 372 (2018), pp. 972–995.
- [6] R. I. SAYE, *High-order quadrature methods for implicitly defined surfaces and volumes in hyperrectangles*, SIAM Journal on Scientific Computing, 37 (2015), pp. A993–A1019.
- [7] M. WICHROWSKI, *A geometric multigrid preconditioner for discontinuous galerkin shifted boundary method*, arXiv preprint arXiv:2506.12899, (2025).
- [8] C.-H. YANG, K. SAURABH, G. SCOVAZZI, C. CANUTO, A. KRISHNAMURTHY, AND B. GANAPATHYSUBRAMANIAN, *Optimal surrogate boundary selection and scalability studies for the shifted boundary method on octree meshes*, Computer Methods in Applied Mechanics and Engineering, 419 (2024), p. 116686.

MW: we have two metrics: estimate of max eigenvalues of Smoother applied to A and number of GMRES iterations